

# Stochastic Multi-Armed Bandits

## Other Algorithms

<sup>1</sup>ECE  
Technical University of Crete

March 10, 2023

# So Far: Explore then Exploit Algorithm

## Definition

Explore-then-Exploit:

- Explore first: Alternate each arm for a total of  $N$  times per arm.
- Exploit after: Pick arm with maximum mean reward (in first  $N$  rounds) - play until round  $T$ .

# So Far: Explore then Exploit Algorithm

## Definition

Explore-then-Exploit:

- Explore first: Alternate each arm for a total of  $N$  times per arm.
- Exploit after: Pick arm with maximum mean reward (in first  $N$  rounds) - play until round  $T$ .

## Performance

- (despite simplicity) can achieve sublinear regret  $O(T^{\frac{2}{3}} \log(T))$
- Key parameter to tune: explore period  $N = O(T^{\frac{2}{3}})$

# So Far: Explore then Exploit Algorithm

## Definition

Explore-then-Exploit:

- Explore first: Alternate each arm for a total of  $N$  times per arm.
- Exploit after: Pick arm with maximum mean reward (in first  $N$  rounds) - play until round  $T$ .

## Performance

- (despite simplicity) can achieve sublinear regret  $O(T^{\frac{2}{3}} \log(T))$
- Key parameter to tune: explore period  $N = O(T^{\frac{2}{3}})$

## What's missing?

- ❶ **Question 1: Can we make the performance “smoother” along the entire  $K$  rounds?** (poor performance in first rounds)
- ❷ **Question 2: Can we do better than  $O(T^{\frac{2}{3}} \log(T))$ ?**

## Definition

$\epsilon$ -Greedy:

- Maintain estimate  $\hat{\mu}_i(t) = \frac{\sum_{n=1}^t r_i^n \cdot X_{i,n}}{\sum_{n=1}^t X_{i,n}}$
- Reminder:  $r_i^n$ : reward of arm  $i$  at round  $n$ ;  $X_{i,n}$ : 1 if arm  $i$  played at round  $n$ .
- At any round  $t$ ,
  - with probability  $\epsilon_t$ : pick a random arm.
  - with probability  $1 - \epsilon_t$ : pick arm  $j : \hat{\mu}_j(t) \geq \hat{\mu}_i(t), \forall i$  (best arm so far)

## Definition

$\epsilon$ -Greedy:

- Maintain estimate  $\hat{\mu}_i(t) = \frac{\sum_{n=1}^t r_i^n \cdot X_{i,n}}{\sum_{n=1}^t X_{i,n}}$
- Reminder:  $r_i^n$ : reward of arm  $i$  at round  $n$ ;  $X_{i,n}$ : 1 if arm  $i$  played at round  $n$ .
- At any round  $t$ ,
  - with probability  $\epsilon_t$ : pick a random arm.
  - with probability  $1 - \epsilon_t$ : pick arm  $j : \hat{\mu}_j(t) \geq \hat{\mu}_i(t), \forall i$  (best arm so far)
- Is this algorithm sublinear also?

## Definition

$\epsilon$ -Greedy:

- Maintain estimate  $\hat{\mu}_i(t) = \frac{\sum_{n=1}^t r_i^n \cdot X_{i,n}}{\sum_{n=1}^t X_{i,n}}$
- Reminder:  $r_i^n$ : reward of arm  $i$  at round  $n$ ;  $X_{i,n}$ : 1 if arm  $i$  played at round  $n$ .
- At any round  $t$ ,
  - with probability  $\epsilon_t$ : pick a random arm.
  - with probability  $1 - \epsilon_t$ : pick arm  $j : \hat{\mu}_j(t) \geq \hat{\mu}_i(t), \forall i$  (best arm so far)
- Is this algorithm sublinear also?
- Is it better than “Explore-then-Exploit”?

# $\epsilon$ -Greedy Algorithm

## Definition

$\epsilon$ -Greedy:

- Maintain estimate  $\hat{\mu}_i(t) = \frac{\sum_{n=1}^t r_i^n \cdot X_{i,n}}{\sum_{n=1}^t X_{i,n}}$
- Reminder:  $r_i^n$ : reward of arm  $i$  at round  $n$ ;  $X_{i,n}$ : 1 if arm  $i$  played at round  $n$ .
- At any round  $t$ ,
  - with probability  $\epsilon_t$ : pick a random arm.
  - with probability  $1 - \epsilon_t$ : pick arm  $j : \hat{\mu}_j(t) \geq \hat{\mu}_i(t), \forall i$  (best arm so far)
- Is this algorithm sublinear also?
- Is it better than “Explore-then-Exploit”?
- Only (key??) parameter to tune is  $\epsilon_t$ ...How?



# $\epsilon$ -Greedy Algorithm

## Definition

$\epsilon$ -Greedy:

- Maintain estimate  $\hat{\mu}_i(t) = \frac{\sum_{n=1}^t r_i^n \cdot X_{i,n}}{\sum_{n=1}^t X_{i,n}}$
- Reminder:  $r_i^n$ : reward of arm  $i$  at round  $n$ ;  $X_{i,n}$ : 1 if arm  $i$  played at round  $n$ .
- At any round  $t$ ,
  - with probability  $\epsilon_t$ : pick a random arm.
  - with probability  $1 - \epsilon_t$ : pick arm  $j : \hat{\mu}_j(t) \geq \hat{\mu}_i(t), \forall i$  (best arm so far)
- Is this algorithm sublinear also?
- Is it better than “Explore-then-Exploit”?
- Only (key??) parameter to tune is  $\epsilon_t$ ...How?

# $\epsilon$ -Greedy vs. Explore-then-Exploit

Question: What are the advantages of  $\epsilon$ -Greedy (if any)?

# $\epsilon$ -Greedy vs. Explore-then-Exploit

Question: What are the advantages of  $\epsilon$ -Greedy (if any)?

- Answer 1: Initial rounds are not *just* exploration  $\rightarrow$  begins to also exploit information immediately.

Question: What are the advantages of  $\epsilon$ -Greedy (if any)?

- Answer 1: Initial rounds are not *just* exploration  $\rightarrow$  begins to also exploit information immediately.
- Answer 2: Never stops exploring  $\rightarrow$  avoids getting stuck with wrong arm.

# How to tune $\epsilon_t$ ?

What about  $\epsilon_t = \epsilon$  (constant)?

- Question: does this have sublinear regret?

# How to tune $\epsilon_t$ ?

## What about $\epsilon_t = \epsilon$ (constant)?

- Question: does this have sublinear regret?
- Answer: No! Easy to prove (will be in homework).

# How to tune $\epsilon_t$ ?

## What about $\epsilon_t = \epsilon$ (constant)?

- Question: does this have sublinear regret?
- Answer: No! Easy to prove (will be in homework).
- $\epsilon_t$  must decrease with  $t$ .

# How to tune $\epsilon_t$ ?

What about  $\epsilon_t = \epsilon$  (constant)?

- Question: does this have sublinear regret?
- Answer: No! Easy to prove (will be in homework).
- $\epsilon_t$  must decrease with  $t$ .

What about  $\epsilon_t = O(t^{-\frac{1}{3}})$ ?



# How to tune $\epsilon_t$ ?

## What about $\epsilon_t = \epsilon$ (constant)?

- Question: does this have sublinear regret?
- Answer: No! Easy to prove (will be in homework).
- $\epsilon_t$  must decrease with  $t$ .

## What about $\epsilon_t = O(t^{-\frac{1}{3}})$ ?

- Question: how many exploration rounds in total?

# How to tune $\epsilon_t$ ?

## What about $\epsilon_t = \epsilon$ (constant)?

- Question: does this have sublinear regret?
- Answer: No! Easy to prove (will be in homework).
- $\epsilon_t$  must decrease with  $t$ .

## What about $\epsilon_t = O(t^{-\frac{1}{3}})$ ?

- Question: how many exploration rounds in total?
- Answer:  $O(T^{-\frac{2}{3}}) \dots$  Why?

# How to tune $\epsilon_t$ ?

## What about $\epsilon_t = \epsilon$ (constant)?

- Question: does this have sublinear regret?
- Answer: No! Easy to prove (will be in homework).
- $\epsilon_t$  must decrease with  $t$ .

## What about $\epsilon_t = O(t^{-\frac{1}{3}})$ ?

- Question: how many exploration rounds in total?
- Answer:  $O(T^{-\frac{2}{3}}) \dots$  Why?
- Question: how many explorations in Explore-then-Exploit?

# How to tune $\epsilon_t$ ?

## What about $\epsilon_t = \epsilon$ (constant)?

- Question: does this have sublinear regret?
- Answer: No! Easy to prove (will be in homework).
- $\epsilon_t$  must decrease with  $t$ .

## What about $\epsilon_t = O(t^{-\frac{1}{3}})$ ?

- Question: how many exploration rounds in total?
- Answer:  $O(T^{-\frac{2}{3}}) \dots$  Why?
- Question: how many explorations in Explore-then-Exploit?
- Answer: Remember,  $N = O(T^{\frac{2}{3}})$  for best sublinear regret.

## Theorem

With  $\epsilon_t = O(t^{-1/3}(k \log t)^{1/3})$ ,  $\epsilon$ -Greedy achieves  $O(t^{2/3}(k \log t)^{1/3})$  regret, for *any*  $t \leq T$

## Theorem

With  $\epsilon_t = O(t^{-1/3}(k \log t)^{1/3})$ ,  $\epsilon$ -Greedy achieves  $O(t^{2/3}(k \log t)^{1/3})$  regret, for *any*  $t \leq T$

## Remarks

- Proof follows similar steps to Explore-then-Exploit.
- Q: Key difference with Explore-then-Exploit?

## Theorem

With  $\epsilon_t = O(t^{-1/3}(k \log t)^{1/3})$ ,  $\epsilon$ -Greedy achieves  $O(t^{2/3}(k \log t)^{1/3})$  regret, for *any*  $t \leq T$

## Remarks

- Proof follows similar steps to Explore-then-Exploit.
- Q: Key difference with Explore-then-Exploit?
- A: The latter has sublinear regret only at the end ( $T$ ).  $\epsilon$ -Greedy has sublinear regret along the way as we wanted (see slide 4)!

## Theorem

With  $\epsilon_t = O(t^{-1/3}(k \log t)^{1/3})$ ,  $\epsilon$ -Greedy achieves  $O(t^{2/3}(k \log t)^{1/3})$  regret, for *any*  $t \leq T$

## Remarks

- Proof follows similar steps to Explore-then-Exploit.
- Q: Key difference with Explore-then-Exploit?
- A: The latter has sublinear regret only at the end ( $T$ ).  $\epsilon$ -Greedy has sublinear regret along the way as we wanted (see slide 4)!
- Q: What is another advantage?



## Theorem

With  $\epsilon_t = O(t^{-1/3}(k \log t)^{1/3})$ ,  $\epsilon$ -Greedy achieves  $O(t^{2/3}(k \log t)^{1/3})$  regret, for *any*  $t \leq T$

## Remarks

- Proof follows similar steps to Explore-then-Exploit.
- Q: Key difference with Explore-then-Exploit?
- A: The latter has sublinear regret only at the end ( $T$ ).  $\epsilon$ -Greedy has sublinear regret along the way as we wanted (see slide 4)!
- Q: What is another advantage?
- A: tuning its key parameter  $\epsilon_t$  does *not* require knowing the horizon!

## Theorem

With  $\epsilon_t = O(t^{-1/3}(k \log t)^{1/3})$ ,  $\epsilon$ -Greedy achieves  $O(t^{2/3}(k \log t)^{1/3})$  regret, for *any*  $t \leq T$

## Remarks

- Proof follows similar steps to Explore-then-Exploit.
- Q: Key difference with Explore-then-Exploit?
- A: The latter has sublinear regret only at the end ( $T$ ).  $\epsilon$ -Greedy has sublinear regret along the way as we wanted (see slide 4)!
- Q: What is another advantage?
- A: tuning its key parameter  $\epsilon_t$  does *not* require knowing the horizon!
- Other sublinear choices for  $\epsilon_t$  also can work (might even do better).

# Can we do even better?

- Q: What is the main disadvantage of *both* algorithms?

# Can we do even better?

- Q: What is the main disadvantage of *both* algorithms?
- A: During exploration, *both* explore *any* arm (even ones that seem bad) with equal probability...

# Can we do even better?

- Q: What is the main disadvantage of *both* algorithms?
- A: During exploration, *both* explore *any* arm (even ones that seem bad) with equal probability...
- Enter “Upper Confidence Bound (UCB)” algorithm.

# Upper Confidence Bound (UCB) Algorithm

## Intuition

Imagine we are at round  $t$ :

- Seems suboptimal to explore arms  $i, j$  with equal probability if  $\hat{\mu}_i(t) \gg \hat{\mu}_j(t)$

# Upper Confidence Bound (UCB) Algorithm

## Intuition

Imagine we are at round  $t$ :

- Seems suboptimal to explore arms  $i, j$  with equal probability if  $\hat{\mu}_i(t) \gg \hat{\mu}_j(t)$
- Goal 1: probability of exploring arm  $i$  at round  $t$  proportional to  $\hat{\mu}_i(t)$

# Upper Confidence Bound (UCB) Algorithm

## Intuition

Imagine we are at round  $t$ :

- Seems suboptimal to explore arms  $i, j$  with equal probability if  $\hat{\mu}_i(t) \gg \hat{\mu}_j(t)$
- Goal 1: probability of exploring arm  $i$  at round  $t$  proportional to  $\hat{\mu}_i(t)$
- Goal 2: BUT...we should still explore an arm  $j$  if it has been explored very few times ( $\rightarrow \text{ourestimate}_j(t)$  is not to be trusted)



# Upper Confidence Bound (UCB) Algorithm

## Intuition

Imagine we are at round  $t$ :

- Seems suboptimal to explore arms  $i, j$  with equal probability if  $\hat{\mu}_i(t) \gg \hat{\mu}_j(t)$
- Goal 1: probability of exploring arm  $i$  at round  $t$  proportional to  $\hat{\mu}_i(t)$
- Goal 2: BUT...we should still explore an arm  $j$  if it has been explored very few times ( $\rightarrow \text{ourestimate}_j(t)$  is not to be trusted)
- Q: How to balance *both* goals?

# Upper Confidence Bound (UCB) Algorithm

## Intuition

Imagine we are at round  $t$ :

- Seems suboptimal to explore arms  $i, j$  with equal probability if  $\hat{\mu}_i(t) \gg \hat{\mu}_j(t)$
- Goal 1: probability of exploring arm  $i$  at round  $t$  proportional to  $\hat{\mu}_i(t)$
- Goal 2: BUT...we should still explore an arm  $j$  if it has been explored very few times ( $\rightarrow \text{ourestimate}_j(t)$  is not to be trusted)
- Q: How to balance *both* goals?
- A: “Optimism under uncertainty”.

# Upper Confidence Bound (UCB) Algorithm

## Definition: Upper Confidence Bound Algorithm

- Maintain estimate  $\hat{\mu}_i(t) = \frac{\sum_{n=1}^t r_i^n \cdot X_{i,n}}{\sum_{n=1}^t X_{i,n}}$
- Maintain estimate  $Q_i(t) = \sum_{n=1}^t X_{i,n}$

# Upper Confidence Bound (UCB) Algorithm

## Definition: Upper Confidence Bound Algorithm

- Maintain estimate  $\hat{\mu}_i(t) = \frac{\sum_{n=1}^t r_i^n \cdot X_{i,n}}{\sum_{n=1}^t X_{i,n}}$
- Maintain estimate  $Q_i(t) = \sum_{n=1}^t X_{i,n}$
- At any round  $t$ :
  - update  $ucb_i(t) = \hat{\mu}_i(t) + \sqrt{\frac{\ln T}{Q_i(t)}}$  (“upper confidence bound”)
  - play arm  $j$ :  $ucb_j(t) \geq ucb_i(t), \forall i$

# Upper Confidence Bound (UCB) Algorithm

## Definition: Upper Confidence Bound Algorithm

- Maintain estimate  $\hat{\mu}_i(t) = \frac{\sum_{n=1}^t r_i^n \cdot X_{i,n}}{\sum_{n=1}^t X_{i,n}}$
- Maintain estimate  $Q_i(t) = \sum_{n=1}^t X_{i,n}$
- At any round  $t$ :
  - update  $ucb_i(t) = \hat{\mu}_i(t) + \sqrt{\frac{\ln T}{Q_i(t)}}$  (“upper confidence bound”)
  - play arm  $j$ :  $ucb_j(t) \geq ucb_i(t), \forall i$
- Arm  $i$  is played either because it has a high predicted reward ( $\hat{\mu}_i(t)$  is large) or because it hasn't been explored enough ( $Q_i(t)$  is small).

# Upper Confidence Bound (UCB) Algorithm

## Definition: Upper Confidence Bound Algorithm

- Maintain estimate  $\hat{\mu}_i(t) = \frac{\sum_{n=1}^t r_i^n \cdot X_{i,n}}{\sum_{n=1}^t X_{i,n}}$
- Maintain estimate  $Q_i(t) = \sum_{n=1}^t X_{i,n}$
- At any round  $t$ :
  - update  $ucb_i(t) = \hat{\mu}_i(t) + \sqrt{\frac{\ln T}{Q_i(t)}}$  (“upper confidence bound”)
  - play arm  $j$ :  $ucb_j(t) \geq ucb_i(t), \forall i$
- Arm  $i$  is played either because it has a high predicted reward ( $\hat{\mu}_i(t)$  is large) or because it hasn't been explored enough ( $Q_i(t)$  is small).
- We will show a simple proof suggesting that UCB achieves better regret  $O(\sqrt{T})$  than previous schemes

# Upper Confidence Bound (UCB) Algorithm

## Definition: Upper Confidence Bound Algorithm

- Maintain estimate  $\hat{\mu}_i(t) = \frac{\sum_{n=1}^t r_i^n \cdot X_{i,n}}{\sum_{n=1}^t X_{i,n}}$
- Maintain estimate  $Q_i(t) = \sum_{n=1}^t X_{i,n}$
- At any round  $t$ :
  - update  $ucb_i(t) = \hat{\mu}_i(t) + \sqrt{\frac{\ln T}{Q_i(t)}}$  (“upper confidence bound”)
  - play arm  $j$ :  $ucb_j(t) \geq ucb_i(t), \forall i$
- Arm  $i$  is played either because it has a high predicted reward ( $\hat{\mu}_i(t)$  is large) or because it hasn't been explored enough ( $Q_i(t)$  is small).
- We will show a simple proof suggesting that UCB achieves better regret  $O(\sqrt{T})$  than previous schemes
- In fact, UCB achieves order optimal regret  $O(\sqrt{T})$  (see Bandit book)

# Upper Confidence Bound (UCB) Algorithm

## Definition: Upper Confidence Bound Algorithm

- Maintain estimate  $\hat{\mu}_i(t) = \frac{\sum_{n=1}^t r_i^n \cdot X_{i,n}}{\sum_{n=1}^t X_{i,n}}$
- Maintain estimate  $Q_i(t) = \sum_{n=1}^t X_{i,n}$
- At any round  $t$ :
  - update  $ucb_i(t) = \hat{\mu}_i(t) + \sqrt{\frac{\ln T}{Q_i(t)}}$  (“upper confidence bound”)
  - play arm  $j$ :  $ucb_j(t) \geq ucb_i(t), \forall i$
- Arm  $i$  is played either because it has a high predicted reward ( $\hat{\mu}_i(t)$  is large) or because it hasn't been explored enough ( $Q_i(t)$  is small).
- We will show a simple proof suggesting that UCB achieves better regret  $O(\sqrt{T})$  than previous schemes
- In fact, UCB achieves order optimal regret  $O()$  (see Bandit book)
- Remark: “order optimal” means that no other algorithm exist that gets better regret in stochastic MAB problems.